



Presto @ Grab



# About the Speaker

Edwin Law - Senior Data Engineering Manager

Been with Grab for 6 years, since August 2013

Managing ~50 Engineers in 4 Data Engineering teams

Reach me at:

email - [edwin.law@grab.com](mailto:edwin.law@grab.com)

Linkedin - [linkedin.com/in/edwinlawhh](https://www.linkedin.com/in/edwinlawhh)

# Grab Today

## DRIVING SOUTHEAST ASIA FORWARD



Grab aims to serve its communities by bringing Southeast Asians closer to what matters to them

Over  
**163 Million**  
mobile downloads



Over  
**9 Million**  
micro-entrepreneurs  
across our network

## 339 CITIES\* 8 COUNTRIES

\*Includes Indonesian cities and regencies, aligning with the government's definition of second-level administrative subdivision under the provinces



Provide the everyday  
services that matter  
the most to our users

Making transportation  
safe, accessible and  
affordable to all

› Seven R&D centres located in Singapore, Beijing, Seattle, Bangalore, Ho Chi Minh City, Kuala Lumpur and Jakarta

# Grab Today

## SERVICES

GRAB OFFERS PAYMENTS AND A RANGE OF CONSUMER SERVICES THROUGH ONE MOBILE APP

### MOBILITY

#### GET A RIDE



##### GrabTaxi

- › E-hailing solves safety and price certainty issues associated with traditional taxis in the region.



##### GrabBike

- › Fastest growing transport service. Most popular option in Greater Jakarta, where population is 60% larger than Beijing.



##### GrabCar

- › Economy and premium options.



##### GrabTukTuk / GrabRemorque / GrabThoneBane

- › Book an affordable ride via our three-wheeler service in the Philippines, Cambodia and Myanmar.



##### GrabWheels

- › A shared active mobility service available in Singapore.

#### GET IT FASTER



##### JustGrab

- › Automatically assigns a vehicle from both Grab's taxi and car fleets at fixed fares for a faster ride.

#### SHARE A RIDE



##### GrabShare

- › Commercial on-demand carpooling service for passengers to share their ride.



##### GrabShuttle

- › Pre-book a shuttle bus seat for an affordable, comfortable commute.

### FINTECH



##### GrabFinancial

- › Micro-lending services to millions of unbanked consumers and micro-entrepreneurs across Southeast Asia.
- › Micro-insurance products for driver-partners to protect their vehicles and livelihoods.



##### GrabPay

- › The only digital payments provider in SEA with access to e-money licenses in the six major ASEAN economies.
- › P2P fund transfers using GrabPay mobile wallet.
- › QR-code enabled payments with merchants in restaurants and shops.



##### GrabRewards

- › Southeast Asia's largest loyalty programme, with over 500 merchants across the region.
- › Users can earn points while using Grab's services and redeem discounts through the GrabRewards catalogue.

### LIFESTYLE



##### GrabExpress

- › On-demand delivery of parcels in Indonesia, Malaysia, Singapore, Thailand and Vietnam.



##### GrabFood

- › An on-demand food delivery platform available in Indonesia, Thailand, Vietnam, Malaysia, Singapore and Philippines.



##### GrabFresh

- › On-demand grocery delivery service available in Indonesia and Thailand.



##### GrabAssistant

- › A concierge service in Philippines that lets customers ask the riders to buy items or queue for them.

# Before Presto, there was ...

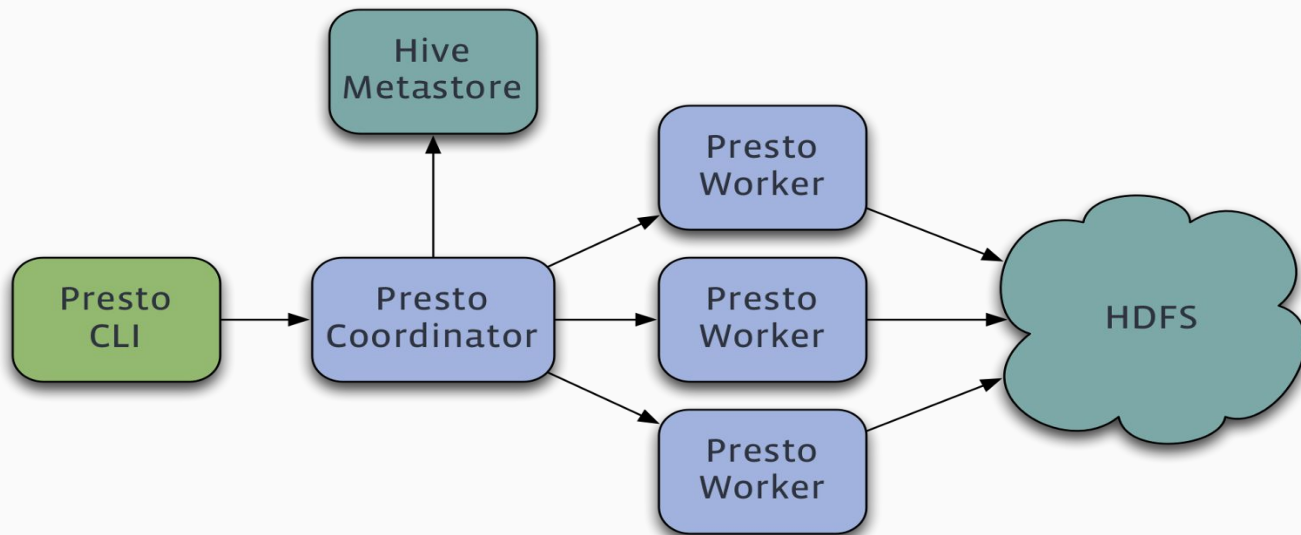


- MPP Data Warehouse , heavily used in 2015 - 2017
- Pros
  - **Access control** features out of the box (Postgres like)
  - **Simple architecture** - one piece of infra to manage
  - storage & compute combined
- Cons
  - **Resource contention** - everyone fighting for same resources
  - Hard (slow) to **scale out** (vastly improved today)
  - **Concurrency limits** (~15 back then, much higher now)

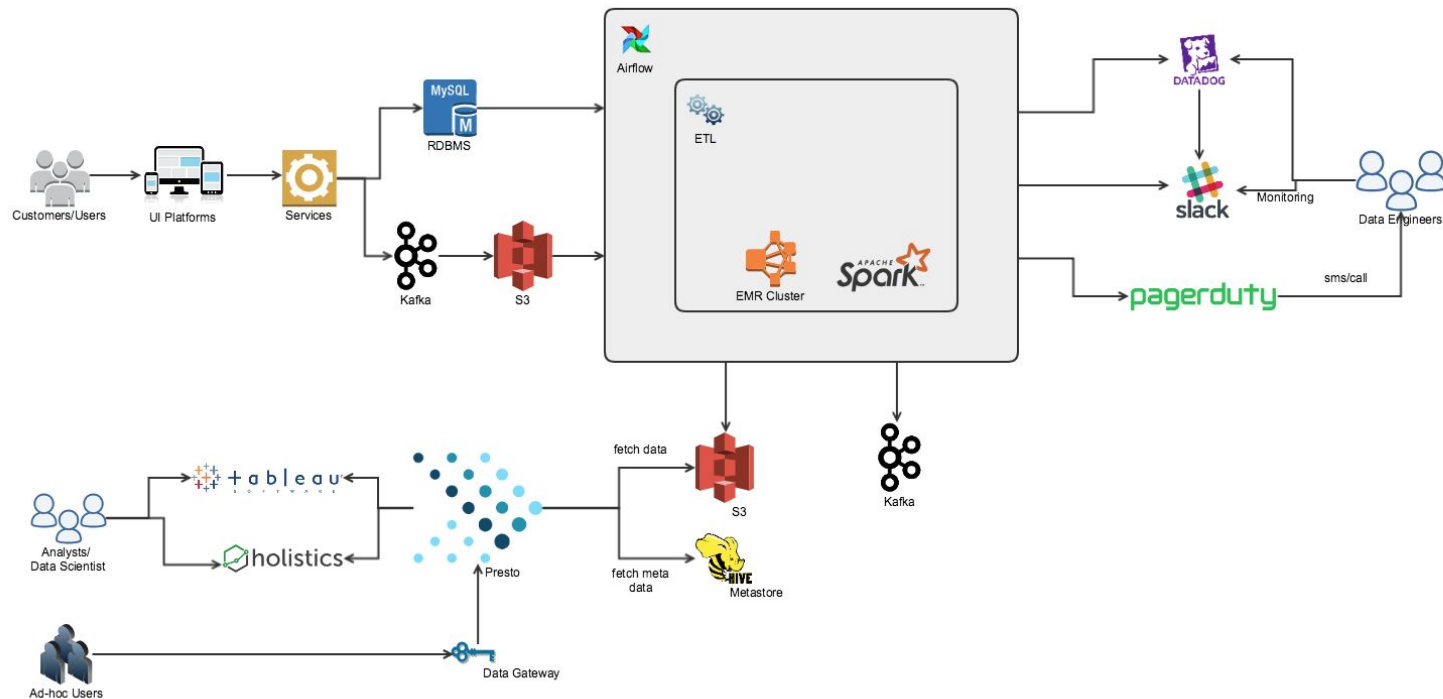


- Our Chosen SQL Query Engine to access our Data Lake on S3
- Pros
  - Enabled **Multi Cluster Architecture** (Decoupled storage and compute)
  - Very fast scale in/out, up/down
  - Great Open Source community
- Cons
  - Complex architecture with many interdependent moving parts
  - Works great but missing many of the database features out of the box
  - Had to build / integrate access control, monitoring, etc.
  - Needs a larger team to support the ecosystem

# Presto Basics



# Our Architecture





# Data Gateway

- Built this early on to get access control for Presto
- Single entry point for all queries
- Authentication and Authorisation service
- Grant access to schemas, tables and clusters
- Integrated with API endpoints of Presto
- Users can route queries to multiple clusters

# Presto on EMR



**amazon**  
**EMR**

- Pros
  - We were already familiar with EMR, use it for our Spark workloads as well
  - Relatively easy to setup, out of the box compatible with our other infrastructure components on AWS
  - Reliable and Scalable
- Cons
  - Configuration Changes were not that easy to apply
  - Cluster Administration interface was clunky
  - No workload aware autoscaling (not simple to setup)
  - High cost due to always on clusters

# Presto on Qubole



- Pros
  - Workload aware Autoscaling Presto Clusters
  - Simple cluster administration, allowed us to
  - Cost saving - 2x workloads at  $\frac{1}{3}$  the original cost on EMR
  - Cross cloud compatibility - AWS & Azure
- Cons
  - Still have workload isolation issues (too many queries from users on some clusters causing long wait times or poor performance)
  - Autoscaling + spot node consumption makes clusters seem less reliable compared to EMR clusters which were always at max size
  - We spun up too many clusters, it became challenging to manage

# Current scale

- 40 workload aware autoscaling Qubole Presto Clusters
- More than 800 nodes at maximum scale
- Master Nodes : r5a.2xlarge (64GB)
- Worker Nodes : r4.8xlarge (> 80% spot nodes) (244GB)
- Serving Adhoc SQL queries and ETL jobs
- 200k+ Queries daily
- 600+ direct users, 3k+ users of BI tools connected to Presto



# Extending the Data Gateway

- The Data Gateway is actually a proxy / query router
- We can make it do more than just access control
- 3 Layer Architecture (Query Router -> Compute Cluster -> Data)
- Extend it to route queries to different clusters based on query & cluster profiles
- Extend it to have it spin up Presto Clusters on K8s when queries come in
- Many other interesting things we can do on the query router layer

# On Demand Presto on Kubernetes



- It takes ~30 seconds to create a new Presto Cluster on Kubernetes
- Every ETL query runs on it's own Presto Cluster
- For Ad-hoc queries, a cluster is created for every user
  - Spin up at first query and torn down when idle for some period

# On Demand Presto on Kubernetes



- **Pros:**

- Ultimate workload isolation
  - each user & each ETL job has dedicated resource
- Best resource utilization through container packing in Kubernetes
- Users don't need to think about resources available, it should always be available!

- **Cons:**

- Added complexity, running Kubernetes is not trivial
- Users might feel like they have less control, they no longer know where their query is running.
- We need to figure out many things on our own, not many ppl doing it like this

# On Demand Presto on Kubernetes



- Allow users to specify what presto version or configuration to use
- Add intelligence that learns based on historical query performance
  - Provision clusters that are just the right size to complete the query in reasonable amount of time
- Dynamically prioritize execution of queries based on query profile / user profile etc.





## What else we're thinking about

- I believe the query router layer allows us to do interesting things.
- Query analysis engine:
  - advise users on how to rewrite their queries
  - Rewrites inefficient user queries before sending it to the compute cluster
- User Submit Standard ANSI SQL - we choose what engine to run it on
  - We could rewrite the query to execute in Presto / Spark / Hive etc.